

# RFR

April 30, 2023

## 1 Random Forest - Regression Task

```
[1]: import numpy as np
import pandas as pd
```

```
[2]: data = pd.read_csv("Boston_Housing.csv")
print(data.head())
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	black	lstat	medv
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

```
[3]: print("The number of observations in the data set = {}".format(data.shape[0]))
```

The number of observations in the data set = 506

```
[6]: X = data.drop('medv', axis=1)
y = data['medv']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[7]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error

param_grid = {
    "n_estimators": [500],
```

```

    "max_features": ["sqrt", "log2"],
    "max_depth": [10, 20, 30],
    "min_samples_split": [2, 4, 6],
    "min_samples_leaf": [1, 2, 4]}

rf = RandomForestRegressor(random_state=42, oob_score=True)

grid_search = GridSearchCV(
    estimator=rf,
    param_grid=param_grid,
    cv=10,
    n_jobs=-1,
    scoring="neg_mean_squared_error")

grid_search.fit(X_train, y_train)

best_rf = grid_search.best_estimator_

print("Best hyperparameters:", grid_search.best_params_)

```

Best hyperparameters: {'max\_depth': 30, 'max\_features': 'sqrt',  
'min\_samples\_leaf': 1, 'min\_samples\_split': 2, 'n\_estimators': 500}

```

[9]: y_pred = best_rf.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
rmse = round(mse ** 0.5, 2)
print("RMSE on testing data:", rmse)

```

RMSE on testing data: 3.2

```

[10]: best_rf.set_params(max_features="sqrt", max_depth=30, min_samples_split=4,
    ↪ min_samples_leaf=1)
best_rf.fit(X_train, y_train)

print(best_rf)

```

RandomForestRegressor(max\_depth=30, max\_features='sqrt', min\_samples\_split=4,  
n\_estimators=500, oob\_score=True, random\_state=42)

```

[11]: print("OOB RMSE:", round(best_rf.oob_score_ ** 0.5, 2))

importance_df = pd.DataFrame({"feature": X_train.columns, "importance": best_rf.
    ↪ feature_importances_})
print(importance_df.sort_values(by="importance", ascending=False))

```

```

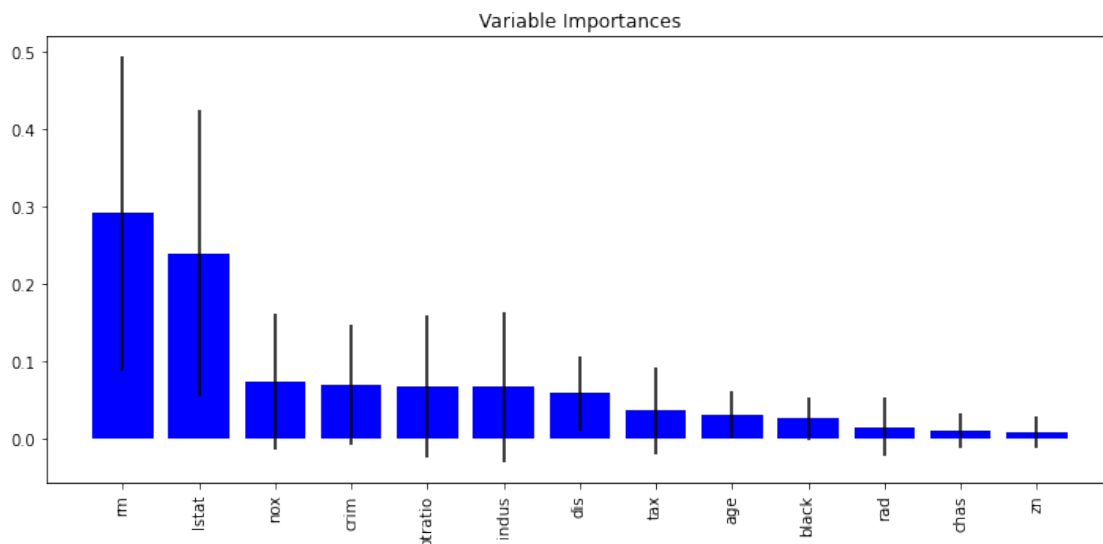
import matplotlib.pyplot as plt
importances = best_rf.feature_importances_
std = np.std([tree.feature_importances_ for tree in best_rf.estimators_],
              ↪axis=0)
indices = np.argsort(importances)[::-1]
features = X_train.columns

plt.figure(figsize=(10, 5))
plt.title("Variable Importances")
plt.bar(range(X_train.shape[1]), importances[indices], color="b",
        ↪yerr=std[indices], align="center")
plt.xticks(range(X_train.shape[1]), features[indices], rotation=90)
plt.xlim([-1, X_train.shape[1]])
plt.tight_layout()
plt.show()

```

OOB RMSE: 0.93

	feature	importance
5	rm	0.291432
12	lstat	0.240054
4	nox	0.073776
0	crim	0.070243
10	ptratio	0.068292
2	indus	0.067541
7	dis	0.058780
9	tax	0.037030
6	age	0.031122
11	black	0.026518
8	rad	0.015804
3	chas	0.010842
1	zn	0.008568



In a regression task using random forest, it is commonly recommended to select approximately the square root of the total number of predictors (variables) at each node split. Therefore for this case it would be either 3 or 4 depending on model performance.