

AMS 580 - Team Project 2 - Python - SVM

Eshan Shakrani, Vishnu Teja Sardee, Priyansh Desai, Mustafa Isik

Importing Necessary Libraries

```
In [ ]: import numpy as np
import pandas as pd

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, accuracy_score
from sklearn.preprocessing import StandardScaler, OneHotEncoder
import matplotlib.pyplot as plt

from sklearn.svm import SVC
```

Problem 1

Read in + Clean the Data

```
In [ ]: # read in the data
df = pd.read_csv('../Data/Titanic.csv')

df.head()
```

Out[]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN

In []: `print("Original Columns:", df.columns.tolist())`

`# remove the "Name", "Ticket", and "Cabin" columns`

`df = df.drop(['PassengerId', 'Name', 'Ticket', 'Cabin'], axis = 1)`

`print("After removing:", df.columns.tolist())`

Original Columns: ['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked']

After removing: ['Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked']

In []: `print(f"There were originally {len(df)} observations.)`

`# remove observations with missing values in the "Age" column`

`df = df.dropna(subset=['Age'])`

`print(f'After removing NA from "Age", there are {len(df)} observations left.)`

There were originally 891 observations.

After removing NA from "Age", there are 714 observations left.

In []: `df.isna().sum()`

```
Out[ ]: Survived      0
         Pclass       0
         Sex         0
         Age         0
         SibSp       0
         Parch       0
         Fare        0
         Embarked    2
         dtype: int64
```

```
In [ ]: # there are still NA values in the "Embarked" column so remove those as well
df = df.dropna()

print(f'There are now {len(df)} observations left.')
```

There are now 712 observations left.

```
In [ ]: stats = df.describe()
stats
```

	Survived	Pclass	Age	SibSp	Parch	Fare
count	712.000000	712.000000	712.000000	712.000000	712.000000	712.000000
mean	0.404494	2.240169	29.642093	0.514045	0.432584	34.567251
std	0.491139	0.836854	14.492933	0.930692	0.854181	52.938648
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	1.000000	20.000000	0.000000	0.000000	8.050000
50%	0.000000	2.000000	28.000000	0.000000	0.000000	15.645850
75%	1.000000	3.000000	38.000000	1.000000	1.000000	33.000000
max	1.000000	3.000000	80.000000	5.000000	6.000000	512.329200

Data Standardization

```
In [ ]: # create a StandardScaler object
ss = StandardScaler()

# fit the scaler to the "Age" and "Fare" columns and transform
df[['Age_Scaled', 'Fare_Scaled']] = ss.fit_transform(df[['Age', 'Fare']])

df.head()
```

Out[]:

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Age_Scaled	Fare_Scaled
0	0	3	male	22.0	1	0	7.2500	S	-0.527669	-0.516380
1	1	1	female	38.0	1	0	71.2833	C	0.577094	0.694046
2	1	3	female	26.0	0	0	7.9250	S	-0.251478	-0.503620
3	1	1	female	35.0	1	0	53.1000	S	0.369951	0.350326
4	0	3	male	35.0	0	0	8.0500	S	0.369951	-0.501257

In []:

```
# remove the original "Age" and "Fare" columns
df = df.drop(['Age', 'Fare'], axis = 1)

df.head()
```

Out[]:

	Survived	Pclass	Sex	SibSp	Parch	Embarked	Age_Scaled	Fare_Scaled
0	0	3	male	1	0	S	-0.527669	-0.516380
1	1	1	female	1	0	C	0.577094	0.694046
2	1	3	female	0	0	S	-0.251478	-0.503620
3	1	1	female	1	0	S	0.369951	0.350326
4	0	3	male	0	0	S	0.369951	-0.501257

In []:

```
print(f'There are currently {len(df)} passengers.')
```

There are currently 712 passengers.

Encoding Categorical Variables

In []:

```
# 'Embarked' and 'Parch' features are both categorical
df['Embarked'].unique()
```

Out[]:

```
array(['S', 'C', 'Q'], dtype=object)
```

In []:

```
df['Sex'].unique()
```

Out[]:

```
array(['male', 'female'], dtype=object)
```

In []:

```
# use a One Hot Encoder to encode the categorical features
ohe = OneHotEncoder(sparse=False)
transformed = ohe.fit_transform(df[['Sex', 'Embarked']])
features_transformed = ohe.get_feature_names_out()
features_transformed
```

Out[]:

```
array(['Sex_female', 'Sex_male', 'Embarked_C', 'Embarked_Q', 'Embarked_S'],
      dtype=object)
```

In []:

```
# add the encoded features to the dataset
df[features_transformed] = transformed
```

In []: df.head()

	Survived	Pclass	Sex	SibSp	Parch	Embarked	Age_Scaled	Fare_Scaled	Sex_female	Sex_m
0	0	3	male	1	0	S	-0.527669	-0.516380	0.0	
1	1	1	female	1	0	C	0.577094	0.694046	1.0	
2	1	3	female	0	0	S	-0.251478	-0.503620	1.0	
3	1	1	female	1	0	S	0.369951	0.350326	1.0	
4	0	3	male	0	0	S	0.369951	-0.501257	0.0	

In []: # drop the original 'Sex' and 'Embarked' columns
df = df.drop(['Sex', 'Embarked'], axis = 1)
df.head()

	Survived	Pclass	SibSp	Parch	Age_Scaled	Fare_Scaled	Sex_female	Sex_male	Embarked_C	E
0	0	3	1	0	-0.527669	-0.516380	0.0	1.0	0.0	
1	1	1	1	0	0.577094	0.694046	1.0	0.0	1.0	
2	1	3	0	0	-0.251478	-0.503620	1.0	0.0	0.0	
3	1	1	1	0	0.369951	0.350326	1.0	0.0	0.0	
4	0	3	0	0	0.369951	-0.501257	0.0	1.0	0.0	

Partitioning the Data

In []: # split the data into features (X) and target (y)
X = df.drop('Survived', axis = 1)
y = df['Survived']

In []: # split the data into 75% training and 25% testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_st

Problem 2

SVM with Linear Kernel

In []: # SVM classifier with Linear kernel and default cost "C"
svc2 = SVC(kernel='linear')

In []: # train the SVC on the training data
svc2.fit(X_train, y_train)

```
Out[ ]: ▾ SVC  
SVC(kernel='linear')
```

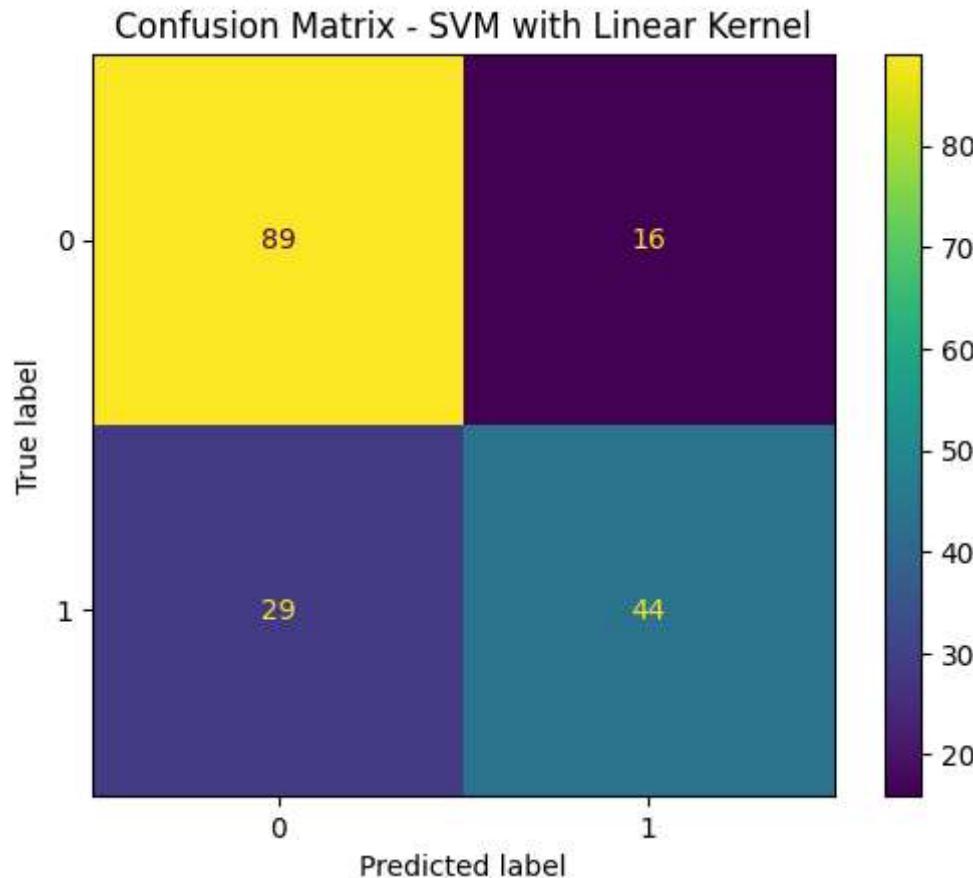
```
In [ ]: # use the SVC to make predictions on the testing data  
y_pred2 = svc2.predict(X_test)
```

Evaluate Performance

```
In [ ]: # generate the confusion matrix  
cm2 = confusion_matrix(y_test, y_pred2, labels = svc2.classes_)  
cm2
```

```
Out[ ]: array([[89, 16],  
               [29, 44]], dtype=int64)
```

```
In [ ]: # better visualization of confusion matrix  
disp_cm2 = ConfusionMatrixDisplay(cm2, display_labels=svc2.classes_)  
disp_cm2.plot()  
plt.title("Confusion Matrix - SVM with Linear Kernel")  
plt.show()
```



```
In [ ]: # extract the elements of the confusion matrix  
tn2, fp2, fn2, tp2 = cm2.ravel()
```

```
In [ ]: # Sensitivity  
sens2 = tp2 / (tp2 + fn2)  
print(f'The probability that a passenger who survived is predicted to have survived
```

The probability that a passenger who survived is predicted to have survived is 0.6027

```
In [ ]: # Specificity  
spec2 = tn2 / (tn2 + fp2)  
print(f'The probability that a passenger who did not survive is predicted to not ha
```

The probability that a passenger who did not survive is predicted to not have survived is 0.8476

```
In [ ]: # Accuracy  
acc2 = sum([tp2, tn2]) / sum([tn2, tp2, fn2, fp2])  
print(f'The overall accuracy of the model is {acc2:.4f}')
```

The overall accuracy of the model is 0.7472

NOTE: Adding the predictions for each model to the testing data at the end

Problem 3

SVM with Linear Kernel

```
In [ ]: # SVM classifier with Linear kernel  
svc3 = SVC(kernel='linear')
```

```
In [ ]: # set up the Grid Search parameters  
param_grid3 = {  
    "C": np.linspace(0, 2, 20)  
}
```

```
In [ ]: # perform the grid search  
grid_search3 = GridSearchCV(svc3, param_grid=param_grid3, cv=5)  
grid_search3.fit(X_train, y_train)
```

```
C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\model_selection\_validation.py:378: FitFailedWarning:
5 fits failed out of a total of 100.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

Below are more details about the failures:

```
5 fits failed with the following error:
Traceback (most recent call last):
  File "C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\mode
l_selection\_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\svm
\_base.py", line 251, in fit
    fit(X, y, sample_weight, solver_type, kernel, random_seed=seed)
  File "C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\svm
\_base.py", line 333, in _dense_fit
    ) = libsvm.fit(
  File "sklearn\svm\_libsvm.pyx", line 192, in sklearn.svm._libsvm.fit
ValueError: C <= 0

    warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\model_selection\_search.py:953: UserWarning: One or more of the test scores are non-finite: [nan 0.7902839 0.7902839 0.7902839 0.7902839 0.7902839 0.7902839
 0.7902839 0.7902839 0.7902839 0.7902839 0.7902839 0.7902839
 0.7902839 0.7902839 0.7902839 0.7902839 0.7902839]
```

Out[]:

```
► GridSearchCV
  ► estimator: SVC
    ► SVC
```

In []: # get the optimal value of C from the GridSearch

```
optimal_C3 = grid_search3.best_params_["C"]
print(f"Optimal value for C: {optimal_C3:.4f}")
```

Optimal value for C: 0.1053

In []: # make the SVM classifier with Linear Kernel and the optimal C value

```
svc3_opt = SVC(kernel="linear", C=optimal_C3)
```

In []: # train the SVC on the training data

```
svc3_opt.fit(X_train, y_train)
```

Out[]:

```
SVC(C=0.10526315789473684, kernel='linear')
```

```
In [ ]: # use the SVC to make predictions on the testing data
y_pred3 = svc3_opt.predict(X_test)
```

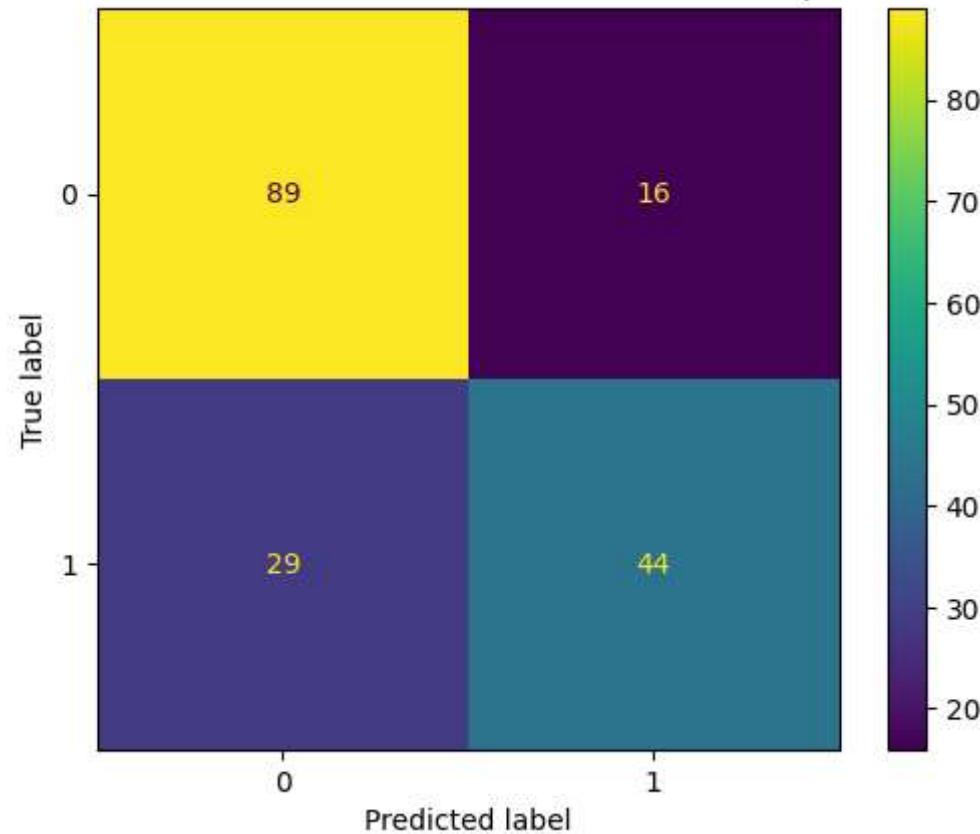
Evaluate Performance

```
In [ ]: # generate the confusion matrix
cm3 = confusion_matrix(y_test, y_pred3, labels = svc3_opt.classes_)
cm3
```

```
Out[ ]: array([[89, 16],
               [29, 44]], dtype=int64)
```

```
In [ ]: # better visualization of confusion matrix
disp_cm3 = ConfusionMatrixDisplay(cm3, display_labels=svc3_opt.classes_)
disp_cm3.plot()
plt.title("Confusion Matrix - SVM with Linear Kernel and Optimal C")
plt.show()
```

Confusion Matrix - SVM with Linear Kernel and Optimal C



```
In [ ]: # extract the elements of the confusion matrix
tn3, fp3, fn3, tp3 = cm3.ravel()
```

```
In [ ]: # Sensitivity
sens3 = tp3 / (tp3 + fn3)
print(f'The probability that a passenger who survived is predicted to have survived
```

The probability that a passenger who survived is predicted to have survived is 0.6027

```
In [ ]: # Specificity  
spec3 = tn3 / (tn3 + fp3)  
print(f'The probability that a passenger who did not survive is predicted to not ha
```

The probability that a passenger who did not survive is predicted to not have survived is 0.8476

```
In [ ]: # Accuracy  
acc3 = sum([tp3, tn3]) / sum([tn3, tp3, fn3, fp3])  
print(f'The overall accuracy of the model is {acc3:.4f}')
```

The overall accuracy of the model is 0.7472

Problem 4

SVM with Radial Basis Kernel

```
In [ ]: # SVM classifier with radial basis kernel  
svc4 = SVC(kernel='rbf')
```

```
In [ ]: # set up the Grid Search parameters  
param_grid4 = {  
    'C': np.linspace(0, 2, 20),  
    'gamma': np.linspace(0, 5, 20)  
}
```

```
In [ ]: # perform the grid search  
grid_search4 = GridSearchCV(svc4, param_grid=param_grid4, cv=5)  
grid_search4.fit(X_train, y_train)
```

```
C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\model_selection\_validation.py:378: FitFailedWarning:  
195 fits failed out of a total of 2000.  
The score on these train-test partitions for these parameters will be set to nan.  
If these failures are not expected, you can try to debug them by setting error_score='raise'.  
  
Below are more details about the failures:  
-----  
100 fits failed with the following error:  
Traceback (most recent call last):  
  File "C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\mode l_selection\_validation.py", line 686, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\svm \_base.py", line 237, in fit  
    raise ValueError(msg)  
ValueError: gamma value must be > 0; 0.0 is invalid. Use a positive number or use 'auto' to set gamma to a value of 1 / n_features.  
  
-----  
95 fits failed with the following error:  
Traceback (most recent call last):  
  File "C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\mode l_selection\_validation.py", line 686, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\svm \_base.py", line 251, in fit  
    fit(X, y, sample_weight, solver_type, kernel, random_seed=seed)  
  File "C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\svm \_base.py", line 333, in _dense_fit  
    ) = libsvm.fit()  
  File "sklearn\svm\_libsvm.pyx", line 192, in sklearn.svm._libsvm.fit  
ValueError: C <= 0  
  
    warnings.warn(some_fits_failed_message, FitFailedWarning)  
C:\Users\eshak\AppData\Roaming\Python\Python310\site-packages\sklearn\model_selection\_search.py:953: UserWarning: One or more of the test scores are non-finite: [  
nan      nan      nan      nan      nan      nan  
      nan      nan      nan      nan      nan      nan  
      nan      nan      nan      nan      nan      nan  
      nan      nan      nan  0.81837418  0.81463587  0.77346147  
  0.69100688  0.63105272  0.61611709  0.61424793  0.61424793  0.61236114  
  0.60675366  0.60299771  0.59924176  0.59924176  0.59924176  0.5973726  
  0.5973726  0.5973726  0.5973726  0.5973726          nan  0.82398166  
  0.81839182  0.8221125  0.81648739  0.77718215  0.74348439  0.69663199  
  0.67603597  0.66101217  0.65353553  0.6460589  0.63667783  0.6292012  
  0.62170693  0.61796861  0.61609945  0.61609945  0.61609945  0.61423029  
          nan  0.82399929  0.82775525  0.8221125  0.82024334  0.81087992  
  0.80333275  0.77152178  0.74528302  0.72659143  0.70784694  0.70784694  
  0.69474519  0.67973902  0.67785223  0.66476812  0.65729148  0.65168401  
  0.65355317  0.65729148          nan  0.82586845  0.82585082  0.82398166  
  0.82398166  0.82398166  0.81461823  0.80710633  0.79770764  0.7808323  
  0.75464645  0.74902134  0.73593722  0.72093105  0.70973373  0.70786457  
  0.70973373  0.70596015  0.70035267  0.69850115          nan  0.8296244  
  0.82771998  0.82398166  0.82585082  0.82022571  0.81461823  0.80338565
```

```

0.79217069 0.782772 0.78275436 0.77150414 0.75842003 0.75279492
0.74531829 0.73969318 0.7340857 0.7247399 0.72100159 0.71350732
    nan 0.83336272 0.82585082 0.81837418 0.82209487 0.82022571
0.81274907 0.80151649 0.79215306 0.7902839 0.78845001 0.78846764
0.77719979 0.7734262 0.75658614 0.75468171 0.73972844 0.72851349
0.7191677 0.71542938      nan 0.83336272 0.83147593 0.82213014
0.82207724 0.81835655 0.81461823 0.79962969 0.79964733 0.79776054
0.79030153 0.77721742 0.77908658 0.77349674 0.76035972 0.75473461
0.74537119 0.74161524 0.73410333 0.73410333      nan 0.83336272
0.82960677 0.82026098 0.82022571 0.82022571 0.81086228 0.80336801
0.80151649 0.80153412 0.79403985 0.79030153 0.7828249 0.76413331
0.76039499 0.76039499 0.75478752 0.74918004 0.73978134 0.7341386
    nan 0.83521425 0.82960677 0.82024334 0.82020808 0.81833892
0.81646976 0.81086228 0.80338565 0.80153412 0.79592664 0.79030153
0.78469406 0.78095574 0.77160995 0.7734791 0.76039499 0.75476988
0.75476988 0.74729325      nan 0.83708341 0.82960677 0.82209487
0.81833892 0.81833892 0.8146006 0.81273144 0.80338565 0.80153412
0.80153412 0.79405749 0.78845001 0.7828249 0.78469406 0.77908658
0.77159231 0.76785399 0.7585082 0.75290072      nan 0.83521425
0.82960677 0.82209487 0.81646976 0.8146006 0.81835655 0.81273144
0.80525481 0.80153412 0.80153412 0.79592664 0.78658085 0.78658085
0.78656322 0.78469406 0.77346147 0.7585082 0.75663904 0.75103156
    nan 0.83708341 0.82960677 0.82022571 0.81833892 0.81646976
0.81835655 0.81273144 0.80338565 0.80153412 0.80153412 0.79592664
0.78845001 0.78845001 0.78471169 0.78469406 0.76787163 0.76039499
0.75478752 0.75480515      nan 0.83708341 0.82960677 0.82020808
0.81833892 0.8146006 0.81835655 0.81273144 0.80338565 0.80153412
0.7977958 0.79592664 0.79031917 0.78845001 0.7828249 0.78469406
0.76787163 0.76041263 0.75480515 0.75103156      nan 0.83708341
0.82773761 0.82207724 0.82020808 0.81648739 0.82022571 0.81273144
0.80338565 0.80153412 0.7977958 0.7977958 0.79031917 0.78658085
0.7828249 0.78658085 0.76788926 0.75667431 0.75291836 0.75103156
    nan 0.83519661 0.82398166 0.82207724 0.82020808 0.81648739
0.82022571 0.81086228 0.80338565 0.79966496 0.7977958 0.79592664
0.79031917 0.78658085 0.78471169 0.78471169 0.76041263 0.75667431
0.75291836 0.7510492      nan 0.83334509 0.82398166 0.82019044
0.82207724 0.81835655 0.81833892 0.81086228 0.80338565 0.79966496
0.7977958 0.79592664 0.78845001 0.78471169 0.78284253 0.78097337
0.75852583 0.75291836 0.75291836 0.75291836      nan 0.83521425
0.82209487 0.81645212 0.82207724 0.82022571 0.81833892 0.81086228
0.80338565 0.79966496 0.7977958 0.79592664 0.7903368 0.78471169
0.78097337 0.77160995 0.75665667 0.75478752 0.75291836 0.7510492
    nan 0.83147593 0.82209487 0.82207724 0.82020808 0.82022571
0.81646976 0.81086228 0.80151649 0.79966496 0.7977958 0.79781344
0.7903368 0.78471169 0.77534826 0.76787163 0.75665667 0.75478752
0.75478752 0.74729325      nan 0.82960677 0.82209487 0.82020808
0.82020808 0.82022571 0.81833892 0.81086228 0.79962969 0.79966496
0.7977958 0.79407512 0.78845001 0.78095574 0.7734791 0.76226415
0.75665667 0.75665667 0.75291836 0.74542409]
warnings.warn(

```

```
Out[ ]: GridSearchCV
         estimator: SVC
             SVC
```

```
In [ ]: # get the optimal values of C and sigma from the GridSearch
optimal_C4, optimal_sigma4 = grid_search4.best_params_["C"], grid_search4.best_params_["sigma"]

print(f'Optimal value for C {optimal_C4:.4f}')
print(f'Optimal value for sigma: {optimal_sigma4:.4f}')
```

Optimal value for C 1.053
Optimal value for sigma: 0.2632

```
In [ ]: # make the SVM classifier with Radial Basis Kernel and the optimal values for C and sigma
svc4_opt = SVC(kernel="rbf", C=optimal_C4, gamma=optimal_sigma4)
```

```
In [ ]: # train the SVC on the training data
svc4_opt.fit(X_train, y_train)
```

```
Out[ ]: SVC
SVC(C=1.0526315789473684, gamma=0.2631578947368421)
```

```
In [ ]: # use the SVC to make predictions on the testing data
y_pred4 = svc4_opt.predict(X_test)
```

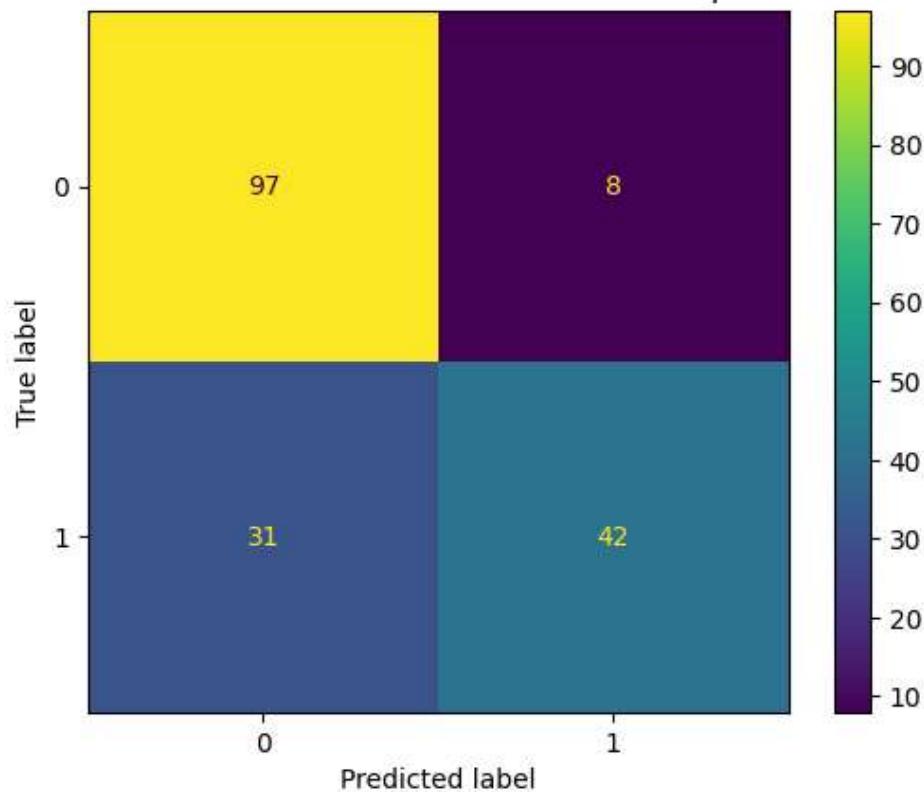
Evaluate Performance

```
In [ ]: # generate the confusion matrix
cm4 = confusion_matrix(y_test, y_pred4, labels = svc4_opt.classes_)
cm4
```

```
Out[ ]: array([[97,  8],
               [31, 42]], dtype=int64)
```

```
In [ ]: # better visualization of confusion matrix
disp_cm4 = ConfusionMatrixDisplay(cm4, display_labels=svc4_opt.classes_)
disp_cm4.plot()
plt.title("Confusion Matrix - SVM with Radial Basis Kernel and Optimal C and Sigma")
plt.show()
```

Confusion Matrix - SVM with Radial Basis Kernel and Optimal C and Sigma



```
In [ ]: # extract the elements of the confusion matrix
tn4, fp4, fn4, tp4 = cm4.ravel()
```

```
In [ ]: # Sensitivity
sens4 = tp4 / (tp4 + fn4)
print(f'The probability that a passenger who survived is predicted to have survived
```

The probability that a passenger who survived is predicted to have survived is 0.5753

```
In [ ]: # Specificity
spec4 = tn4 / (tn4 + fp4)
print(f'The probability that a passenger who did not survive is predicted to not ha
```

The probability that a passenger who did not survive is predicted to not have survived is 0.9238

```
In [ ]: # Accuracy
acc4 = sum([tp4, tn4]) / sum([tn4, tp4, fn4, fp4])
print(f'The overall accuracy of the model is {acc4:.4}')
```

The overall accuracy of the model is 0.7809

Problem 5

SVM with Polynomial Kernel

```
In [ ]: # SVM classifier with polynomial kernel
svc5 = SVC(kernel='poly')
```

```
In [ ]: # set up the Grid Search parameters
param_grid5 = {
    "C": np.linspace(0, 2, 20),
    'degree': list(range(1,6)),
    'coef0': [-1, 0, 1]
}
```

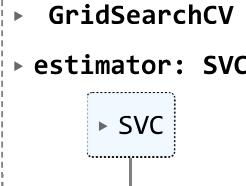
```
In [ ]: # perform the grid search
grid_search5 = GridSearchCV(svc5, param_grid=param_grid5, cv=5)
grid_search5.fit(X_train, y_train)
```



```
0.7902839 0.38576971 0.53554929 0.42322342 0.51126785 0.7902839
0.81650503 0.82960677 0.80349145 0.79971786 0.7902839 0.82024334
0.83147593 0.81840945 0.80525481 0.7902839 0.38576971 0.53366249
0.42135426 0.51126785 0.7902839 0.81463587 0.82586845 0.80911656
0.80532534 0.7902839 0.82024334 0.82773761 0.81467113 0.80525481
0.7902839 0.38763886 0.53369776 0.41387762 0.51500617 0.7902839
0.81650503 0.82773761 0.8222183 0.80158702 0.7902839 0.81835655
0.82773761 0.81467113 0.80338565 0.7902839 0.38576971 0.52995944
0.41387762 0.51126785 0.7902839 0.81650503 0.82773761 0.82220067
0.80158702 0.7902839 0.82024334 0.82773761 0.8146535 0.79962969
0.7902839 0.38390055 0.52809028 0.41012167 0.51313701 0.7902839
0.81461823 0.82960677 0.81657556 0.80160466 0.7902839 0.82024334
0.82960677 0.81278434 0.79962969 0.7902839 0.38390055 0.52622113
0.40825251 0.51126785 0.7902839 0.81648739 0.82960677 0.81657556
0.81098572 0.7902839 0.82024334 0.82773761 0.81278434 0.80148122
0.7902839 0.38201375 0.52622113 0.40638335 0.5093987 0.7902839
0.81648739 0.82773761 0.81844472 0.81285488 0.7902839 0.82024334
0.82773761 0.8146535 0.80336801 0.7902839 0.38390055 0.52246517
0.40638335 0.51126785 0.7902839 0.81648739 0.82960677 0.8221654
0.81472403 0.7902839 0.82213014 0.82773761 0.81278434 0.80149885]
```

warnings.warn(

Out[]:



In []:

```
# get the optimal values of C, degree, and scale from the GridSearch
optimal_C5, optimal_degree5, optimal_scale5 = grid_search5.best_params_["C"], grid_
print(f'Optimal value for C {optimal_C5:.4f}')
print(f'Optimal value for degree: {optimal_degree5}')
print(f'Optimal value for scale: {optimal_scale5}')
```

Optimal value for C 1.158

Optimal value for degree: 3

Optimal value for scale: 0

In []:

```
# make the SVM classifier with Polynomial Kernel and the optimal values for C, degr
svc5_opt = SVC(kernel="poly", C=optimal_C5, degree=optimal_degree5, coef0=optimal_s
```

In []:

```
# train the SVC on the training data
svc5_opt.fit(X_train, y_train)
```

Out[]:

```
SVC
SVC(C=1.1578947368421053, coef0=0, kernel='poly')
```

In []:

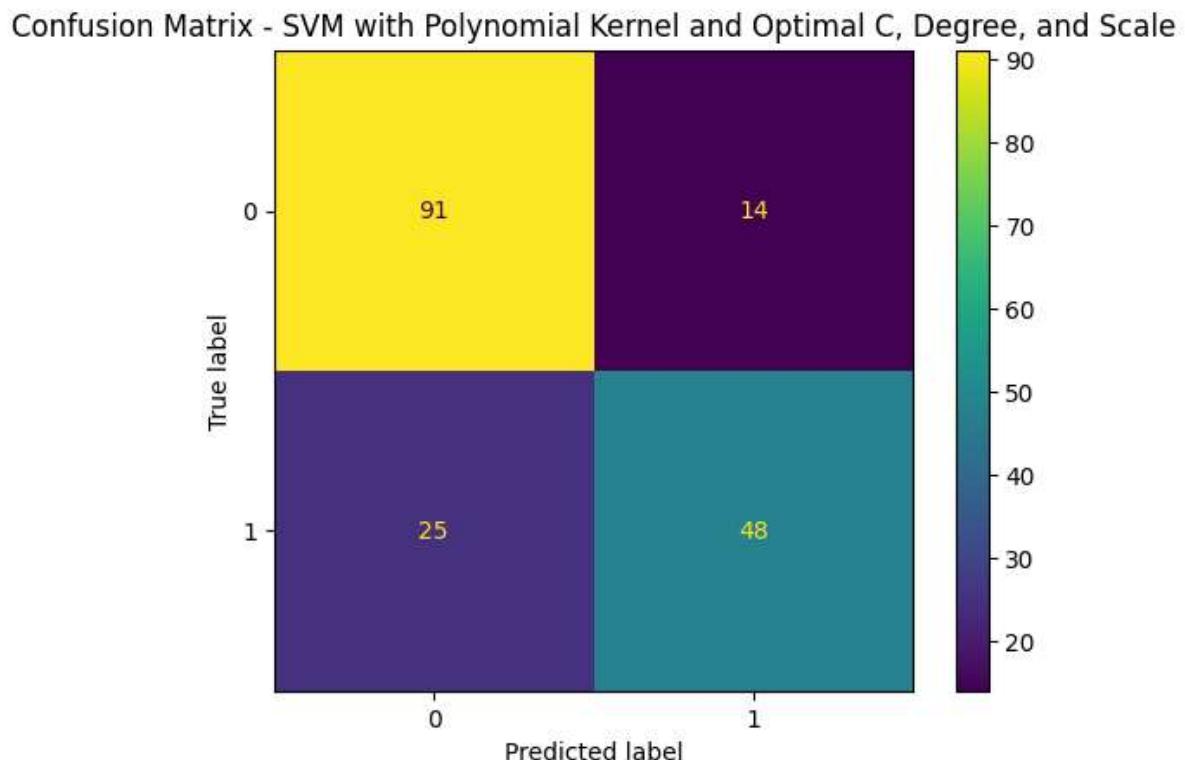
```
# use the SVC to make predictions on the testing data
y_pred5 = svc5_opt.predict(X_test)
```

Evaluate Performance

```
In [ ]: # generate the confusion matrix
cm5 = confusion_matrix(y_test, y_pred5, labels = svc5_opt.classes_)
cm5
```

```
Out[ ]: array([[91, 14],
 [25, 48]], dtype=int64)
```

```
In [ ]: # better visualization of confusion matrix
disp_cm5 = ConfusionMatrixDisplay(cm5, display_labels=svc5_opt.classes_)
disp_cm5.plot()
plt.title("Confusion Matrix - SVM with Polynomial Kernel and Optimal C, Degree, and Scale")
plt.show()
```



```
In [ ]: # extract the elements of the confusion matrix
tn5, fp5, fn5, tp5 = cm5.ravel()
```

```
In [ ]: # Sensitivity
sens5 = tp5 / (tp5 + fn5)
print(f'The probability that a passenger who survived is predicted to have survived
```

The probability that a passenger who survived is predicted to have survived is 0.6575

```
In [ ]: # Specificity
spec5 = tn5 / (tn5 + fp5)
print(f'The probability that a passenger who did not survive is predicted to not ha
```

The probability that a passenger who did not survive is predicted to not have survived is 0.8667

```
In [ ]: # Accuracy
acc5 = sum([tp5, tn5]) / sum([tn5, tp5, fn5, fp5])
```

```
print(f'The overall accuracy of the model is {acc5:.4f}')
```

The overall accuracy of the model is 0.7809

Adding the predictions from Problem 2 to the testing data

```
In [ ]: y_test_df = pd.DataFrame(y_test)
```

```
In [ ]: y_test_df['Problem 2 Predictions'] = y_pred2
```

```
In [ ]: y_test_df
```

```
Out[ ]:   Survived  Problem 2 Predictions
```

	Survived	Problem 2 Predictions
509	1	0
640	0	0
537	1	1
526	1	1
262	0	0
...
741	0	0
562	0	0
586	0	0
88	1	1
187	1	0

178 rows × 2 columns

Problem 6

Find the model with the best accuracy

```
In [ ]: # acc3: Linear with optimal C
# acc4: Radial with optimal C and sigma
# acc5: Polynomial with optimal C, degree, and scale
accuracies = [acc3, acc4, acc5]

models = ['Linear SVM', 'SVM with Radial Basis Kernel', 'SVM with Polynomial Kernel']

# get the highest accuracy of the three
max_accuracy = max(accuracies)

# get the highest accuracy's index
max_index = accuracies.index(max_accuracy)

# get the highest accuracy's corresponding model
best_model = models[max_index]
```

```
print(f'Best model was {best_model} with an accuracy of {max_accuracy:.4}')
```

Best model was SVM with Radial Basis Kernel with an accuracy of 0.7809