

# Neural Networks, Regression Task, Exercise.

name

xx/xx/xxxx

1. Divide the data into 75% training and 25% testing & normalize the data

```
if (!requireNamespace("tidyverse")) install.packages('tidyverse')

## Loading required namespace: tidyverse

if (!requireNamespace("caret")) install.packages('caret')

## Loading required namespace: caret

if (!requireNamespace("neuralnet")) install.packages('neuralnet')

## Loading required namespace: neuralnet

library(tidyverse)

## -- Attaching packages ----- tidyverse 1.3.1 --

## v ggplot2 3.3.5      v purrr  0.3.4
## v tibble  3.1.6      v dplyr  1.0.8
## v tidyr   1.2.0      v stringr 1.4.0
## v readr   2.0.1      v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(caret)

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
## lift
```

```
library(neuralnet)
```

```
##  
## Attaching package: 'neuralnet'  
  
## The following object is masked from 'package:dplyr':  
##  
##      compute
```

```
library(MASS)
```

```
##  
## Attaching package: 'MASS'  
  
## The following object is masked from 'package:dplyr':  
##  
##      select
```

```
data("Boston")  
data = Boston  
data <- subset(data, select = -c(rad))  
# mean & standard deviation of the response  
mean = mean(data$medv)  
sd = sd(data$medv)  
# normalize the data  
data = data.frame(scale(data))  
  
set.seed(123)  
training.samples <- data$medv %>%  
  createDataPartition(p = 0.75, list = FALSE)  
train.data <- data[training.samples, ]  
test.data <- data[-training.samples, ]  
str(train.data) # 381 obs
```

```
## 'data.frame':   381 obs. of  13 variables:  
## $ crim      : num  -0.419 -0.417 -0.416 -0.412 -0.41 ...  
## $ zn        : num   0.2845 -0.4872 -0.4872 -0.4872 0.0487 ...  
## $ indus     : num  -1.287 -0.593 -1.306 -1.306 -0.476 ...  
## $ chas      : num  -0.272 -0.272 -0.272 -0.272 -0.272 ...  
## $ nox       : num  -0.144 -0.74 -0.834 -0.834 -0.265 ...  
## $ rm        : num   0.413 0.194 1.015 1.227 -0.388 ...  
## $ age       : num  -0.1199 0.3668 -0.8091 -0.5107 -0.0702 ...  
## $ dis       : num   0.14 0.557 1.077 1.077 0.838 ...  
## $ tax       : num  -0.666 -0.986 -1.105 -1.105 -0.577 ...  
## $ ptratio   : num  -1.458 -0.303 0.113 0.113 -1.504 ...  
## $ black     : num   0.441 0.441 0.416 0.441 0.426 ...  
## $ lstat     : num  -1.0745 -0.492 -1.3602 -1.0255 -0.0312 ...  
## $ medv      : num   0.1595 -0.1014 1.1816 1.486 0.0399 ...
```

```
str(test.data) # 125 obs
```

```
## 'data.frame': 125 obs. of 13 variables:
## $ crim : num -0.417 -0.417 -0.396 -0.394 -0.347 ...
## $ zn : num -0.4872 -0.4872 0.0487 0.0487 -0.4872 ...
## $ indus : num -0.593 -1.306 -0.476 -0.476 -0.437 ...
## $ chas : num -0.272 -0.272 -0.272 -0.272 -0.272 ...
## $ nox : num -0.74 -0.834 -0.265 -0.265 -0.144 ...
## $ rm : num 1.281 0.207 -0.93 0.131 -0.478 ...
## $ age : num -0.266 -0.351 1.116 0.914 -0.241 ...
## $ dis : num 0.557 1.077 1.086 1.212 0.433 ...
## $ tax : num -0.986 -1.105 -0.577 -0.577 -0.601 ...
## $ ptratio: num -0.303 0.113 -1.504 -1.504 1.175 ...
## $ black : num 0.396 0.41 0.328 0.393 0.441 ...
## $ lstat : num -1.208 -1.042 2.419 1.092 -0.615 ...
## $ medv : num 1.323 0.671 -0.656 -0.819 -0.232 ...
```

2. NN model with (i) no hidden layer, (ii) the default loss function of 'sse', and (iii) the default activation function of 'identity'.

```
set.seed(123)
nn = neuralnet(medv~., data = train.data, hidden = 0, err.fct = "sse", linear.output = T)
plot(nn, rep = 'best')
```

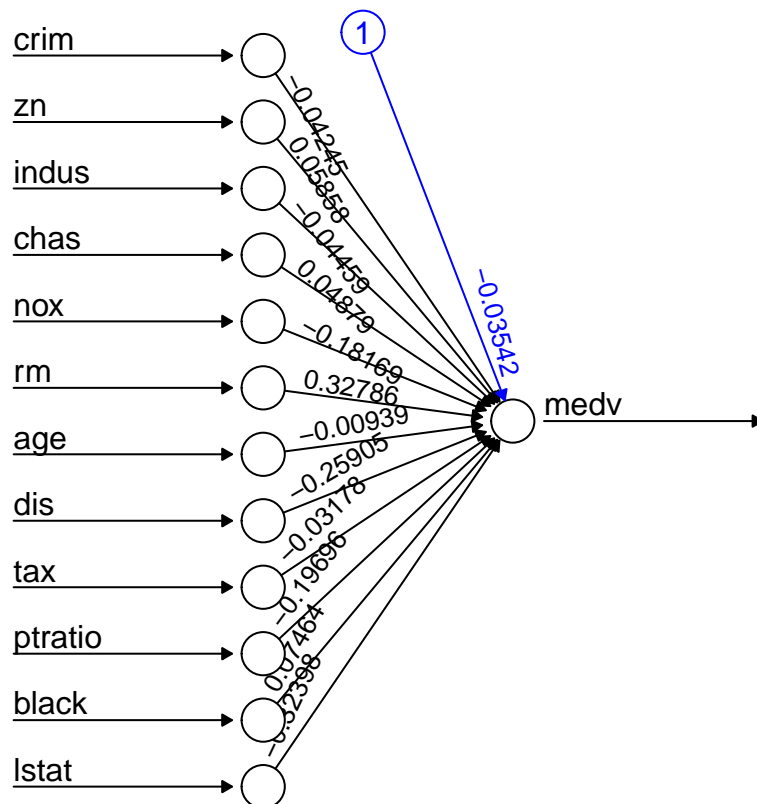


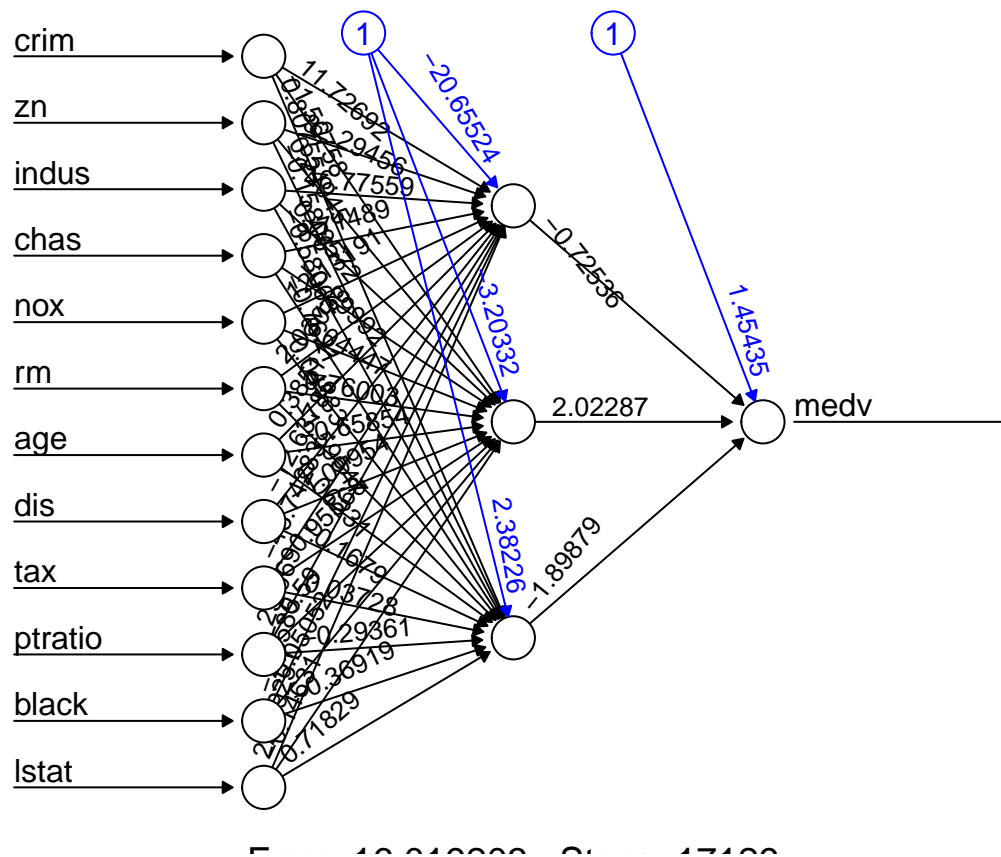
Figure 20: 205234 - Class 100

```
pr.nn0 = predict(nn, test.data)
# Test MSE
(MSE.nn.1 = RMSE(test.data$medv*sd+mean, pr.nn0*sd+mean)^2)
```

```
## [1] 42.90577
```

3. NN model with (i) one hidden layer with 3 neurons, (ii) the default loss function of 'sse', and (iii) the default activation function of 'identity'.

```
set.seed(123)
nn = neuralnet(medv~., data = train.data, hidden = 3, err.fct = "sse", linear.output = T)
plot(nn, rep = 'best')
```



```
pr.nn1 = predict(nn, test.data)
# Test MSE
(MSE.nn.2 = RMSE(test.data$medv*sd+mean, pr.nn1*sd+mean)^2)
```

```
## [1] 37.01437
```

4. MLR model

```
set.seed(123)
mlr = lm(medv~., data = train.data)
summary(mlr)
```

```
##
## Call:
## lm(formula = medv ~ ., data = train.data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.40013 -0.26214 -0.06769  0.17814  2.35738
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.035423   0.023695  -1.495 0.135789
## crim        -0.042440   0.031485  -1.348 0.178506
## zn           0.058543   0.035236   1.661 0.097474 .
## indus       -0.044564   0.046408  -0.960 0.337545
## chas         0.048801   0.026021   1.875 0.061518 .
## nox         -0.181812   0.050075  -3.631 0.000323 ***
## rm           0.327836   0.032518  10.082 < 2e-16 ***
## age         -0.009319   0.042395  -0.220 0.826144
## dis         -0.259040   0.045304  -5.718 2.23e-08 ***
## tax         -0.031723   0.043656  -0.727 0.467900
## ptratio     -0.197012   0.030720  -6.413 4.38e-10 ***
## black        0.074621   0.026127   2.856 0.004533 **
## lstat       -0.324002   0.044322  -7.310 1.68e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4609 on 368 degrees of freedom
## Multiple R-squared:  0.7626, Adjusted R-squared:  0.7549
## F-statistic: 98.52 on 12 and 368 DF,  p-value: < 2.2e-16
```

```
pr.mlr = predict(mlr, test.data)
# Test MSE
(MSE.mlr = RMSE(test.data$medv*sd+mean, pr.mlr*sd+mean)^2)
```

```
## [1] 42.90551
```

```
# Compare MSE
print(paste(MSE.nn.1, MSE.nn.2, MSE.mlr))
```

```
## [1] "42.9057695828218 37.0143749146895 42.9055116896361"
```

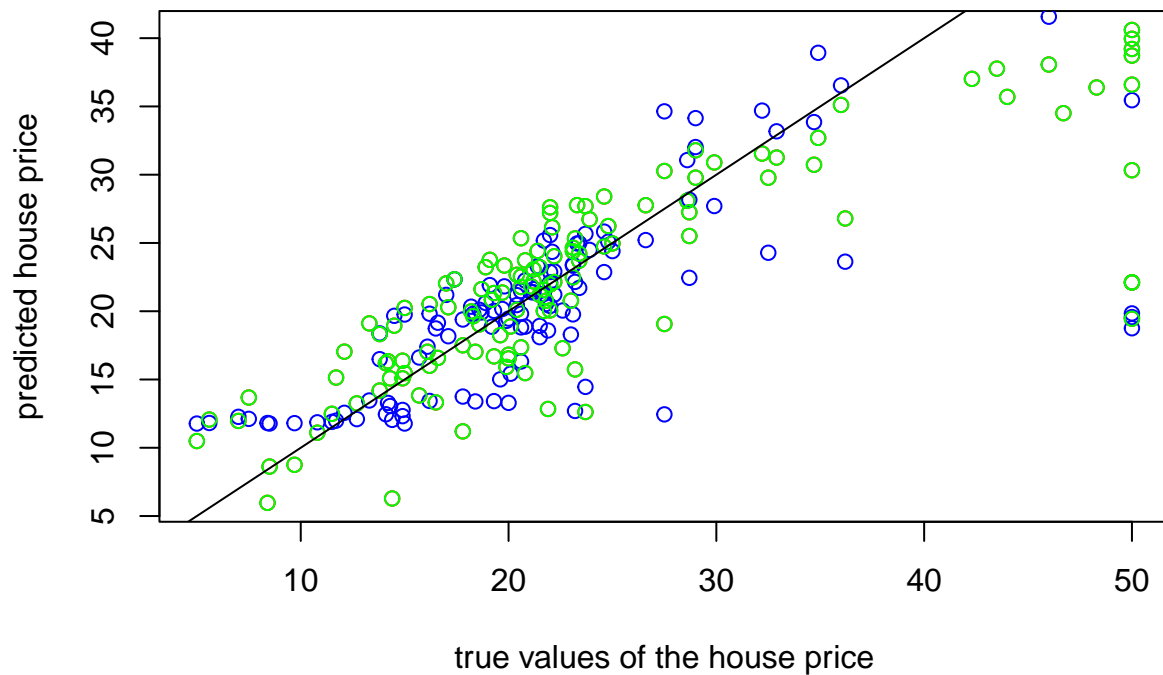
```
# Compare with multiple linear regression
# summarize the predictions from different models
final1 <- data.frame(predictions_NN0=pr.nn0*sd+mean, predictions_NN1=pr.nn1*sd+mean, predictions_MLR =pr
knitr::kable(head(final1))
```

|    | predictions_NN0 | predictions_NN1 | predictions_MLR | actual_response |
|----|-----------------|-----------------|-----------------|-----------------|
| 3  | 30.73145        | 33.85545        | 30.73168        | 34.7            |
| 6  | 25.50940        | 22.44689        | 25.50949        | 28.7            |
| 9  | 13.32310        | 18.73894        | 13.32406        | 16.5            |
| 11 | 20.24198        | 19.76108        | 20.24290        | 15.0            |

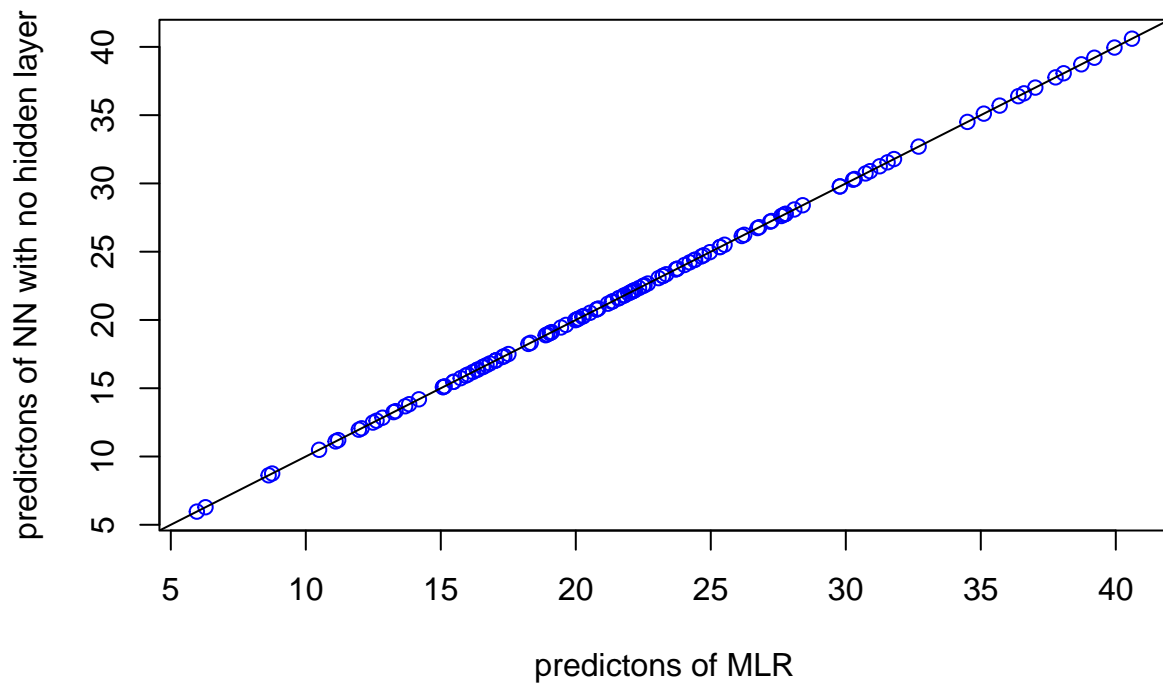
|    | predictions_NN0 | predictions_NN1 | predictions_MLR | actual_response |
|----|-----------------|-----------------|-----------------|-----------------|
| 14 | 20.10770        | 20.45413        | 20.10700        | 20.4            |
| 15 | 19.98503        | 20.33412        | 19.98474        | 18.2            |

```
attach(final1)
```

```
# NN model with no hidden layer, with one hidden layer with 3 neurons and MLR vs. the true values
plot(actual_response, predictions_NN0, col="red", ylab = 'predicted house price', xlab = 'true values of the house price')
points(actual_response, predictions_NN1, col="blue")
points(actual_response, predictions_MLR, col="green")
abline(a = 0, b = 1)
```



```
# NN model with one hidden layer with 3 neurons vs. MLR
plot(predictions_MLR, predictions_NN0, col="blue", ylab = 'predictons of NN with no hidden layer', xlab = 'predictons of NN with one hidden layer')
abline(a = 0, b = 1)
```



5. NN model with (i) one hidden layer with 3 neurons, (ii) the default loss function of “sse”, and (iii) the output layer with the default activation function of ‘identity’, but the hidden layer with the activation function of ‘relu’.

```
# devtools::install_github('rstudio/cloudml')
library(keras)
library(dplyr)
library(cloudml)
```

```
## Loading required package: tfruns
```

```
train_x = subset(train.data, select = -medv)
train_x_s = scale(train_x)
train_y = as.matrix(subset(train.data, select = medv))
test_x = subset(test.data, select = -medv)
test_x_s = scale(test_x)
test_y = as.matrix(subset(test.data, select = medv))

set.seed(123)
model <- keras_model_sequential()
```

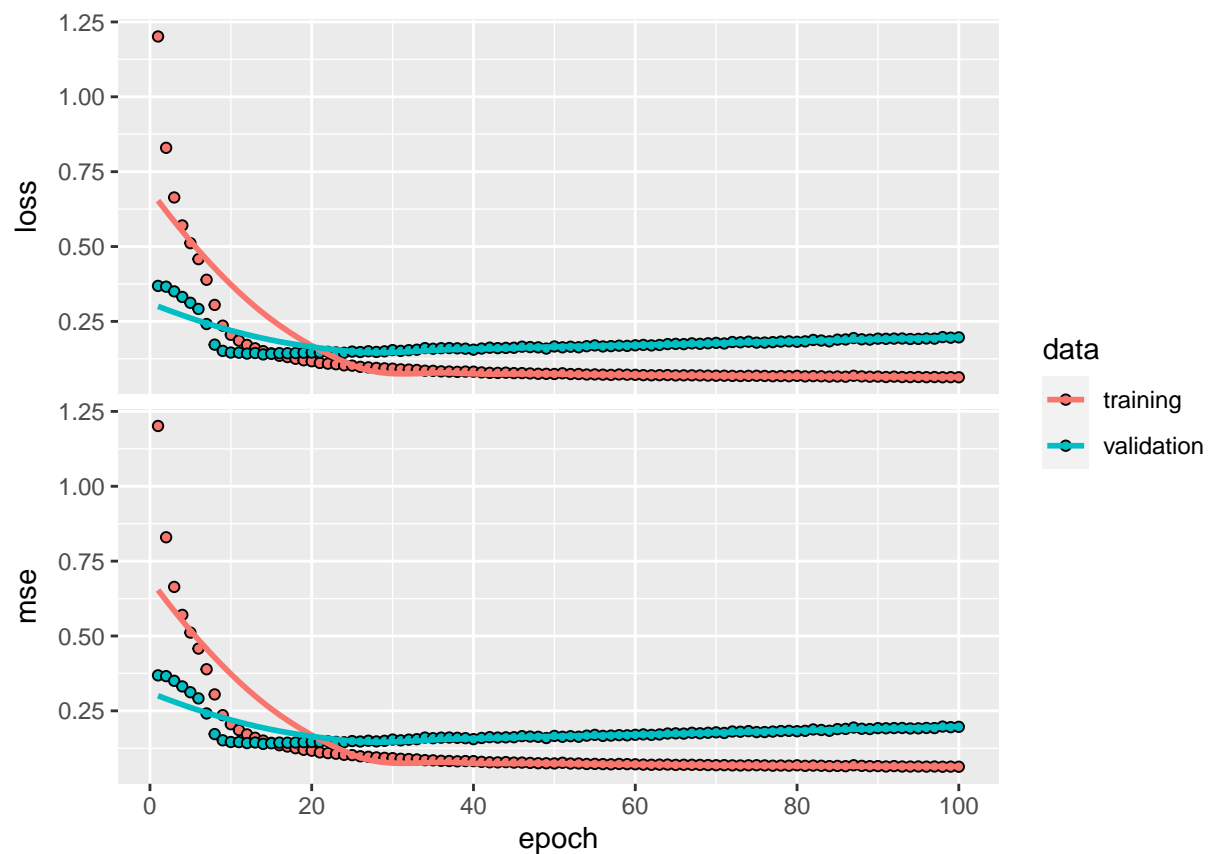
```
## Loaded Tensorflow version 2.4.4
```

```
model %>% layer_dense(units = 12, activation = 'relu', input_shape = c(12)) %>%
  layer_dense(units = 3, activation = "relu") %>%
  layer_dense(units = 1, activation = "linear")
model %>% compile(loss='mse',optimizer='adam',metrics='mse')
summary(model)
```

```
## Model: "sequential"
## -----
## Layer (type)                Output Shape          Param #
## =====
## dense_2 (Dense)              (None, 12)             156
## -----
## dense_1 (Dense)              (None, 3)              39
## -----
## dense (Dense)                (None, 1)              4
## =====
## Total params: 199
## Trainable params: 199
## Non-trainable params: 0
## -----
```

```
history = model %>% fit(train_x_s,train_y, epochs=100,batch_size = 8,validation_split = 0.1)
plot(history)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



```
preds <- predict(model, test_x_s)

# test MSE
RMSE(test.data$medv*sd+mean, preds*sd+mean)^2
```

```
## [1] 16.63245
```



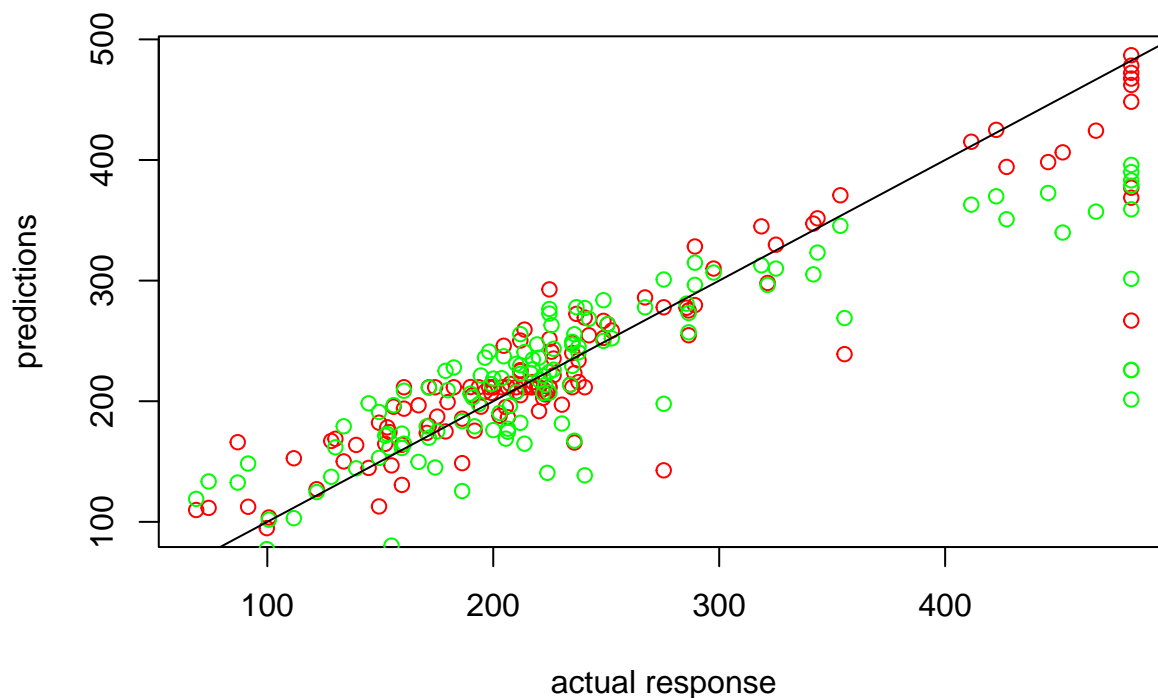
```
# Compare with multiple linear regression
final2 <- data.frame(predictions_NN_RELU=preds*sd+mean, predictions_MLR =pr.mlr*sd+mean, actual_response=
knitr::kable(head(final2))
```

|    | predictions_NN_RELU | predictions_MLR | medv |
|----|---------------------|-----------------|------|
| 3  | 35.31134            | 30.73168        | 34.7 |
| 6  | 25.24495            | 25.50949        | 28.7 |
| 9  | 20.55788            | 13.32406        | 16.5 |
| 11 | 20.55788            | 20.24290        | 15.0 |
| 14 | 20.16028            | 20.10700        | 20.4 |
| 15 | 20.57497            | 19.98474        | 18.2 |

```
attach(final2)
```

```
## The following object is masked from final1:
##
## predictions_MLR
```

```
plot(actual_response*sd+mean, predictions_NN_RELU*sd+mean, col="red", ylab = 'predictions', xlab = 'actual
points(actual_response*sd+mean, predictions_MLR*sd+mean, col="green")
abline(a = 0, b = 1)
```



guidance of how to deal with binary string variables

The variable 'w4' in Quiz 7 is a string variable (Y/N), you can clean it in this way

```
(w4 = c(rep('Y',5), rep('N',5)))
```

```
## [1] "Y" "Y" "Y" "Y" "Y" "N" "N" "N" "N" "N"
```

```
# replace 'Y' with 1 and 'N' with 0:  
library(dplyr)  
(w4 = ifelse(w4 == 'N', 0,1))
```

```
## [1] 1 1 1 1 1 0 0 0 0 0
```

keras example <https://www.datatechnotes.com/2019/01/regression-example-with-keras-in-r.html>

python installation <https://www.dataquest.io/blog/installing-python-on-mac/> <https://phoenixnap.com/kb/how-to-install-python-3-windows>